

# Unconstrained Optimization Assignment

Andrea Cognolato, s281940@studenti.polito.it

January 25, 2021

## 1 Introduction

Descent methods are a vast family of iterative optimization algorithms for finding a local minimum of a differentiable function  $f$ . They operate by repeatedly moving in a direction where  $f$  decreases.

More rigorously, descent methods aim to solve the following unconstrained optimization problem:

$$\min_x f(x)$$

where  $f : \mathbb{R}^n \rightarrow \mathbb{R}$ ,  $x \in \mathbb{R}^n$ .

To do so, a direction  $p_k$  is chosen such that  $\nabla f(x_k)^T p_k < 0$  and the next iterate  $x_{k+1}$  is generated by moving along  $p_k$  with a step of length  $\alpha_k$ . That is:  $x_{k+1} = x_k + \alpha_k p_k$ .

There are several ways to choose  $p_k$ . In this assignment, we will consider three methods: Steepest Descent, Fletcher Reeves, and Polak–Ribière.

Likewise, we have many possible choices for choosing the step length  $\alpha_k$ , which result in various *line search* strategies. In our case we will use backtracking line search, a strategy based on repeatedly shrinking our step length  $\alpha$  until the Armijo conditions are satisfied.

We implement these three methods in Julia<sup>1</sup>, and analyze their computation time on problems of size  $n = 10^4, 10^5$  and using both the exact gradient, forward finite differences, and centered finite differences.

### 1.1 Steepest Descent

Steepest descent (or gradient descent) is one of the simplest methods. At every iteration the steepest direction, i.e. the direction along with  $f$  decreases the most, is picked.

The gradient of a function at a point is denoted by  $\nabla f(x_k)$ , and is a vector indicating the direction of most rapid ascent, by taking its opposite, we obtain the direction of steepest descent.

$$p_k^{SD} = -\nabla f(x_k)$$

Steepest descent unfortunately can display pathological behaviour on functions with very narrow valleys. In this case, the algorithm assumes a “zig-zagging” behaviour where a lots of tiny steps are made towards the solution, drastically decreasing its convergence speed.

### 1.2 Fletcher-Reeves

To solve steepest descent’s the issue with very narrow valleys, nonlinear conjugate gradient methods, such as Fletcher-Reeves were developed. To compute their descent directions, they combine the direction of steepest descent with the descent directions of the previous step, weighted by a factor  $\beta_k$ .

---

<sup>1</sup>Full codes available at [github.com/mrandri19/polito-numerical-optimization](https://github.com/mrandri19/polito-numerical-optimization)

$$d_k = -\nabla f(x_k)$$

$$p_{k+1} = d_k + \beta_k p_k$$

In particular, in Fletcher-Reeves, the coefficient  $\beta_k$  is computed as follows:

$$\beta_k^{FR} = \frac{\nabla f(x_{k+1})^T \nabla f(x_{k+1})}{\nabla f(x_k)^T \nabla f(x_k)}$$

### 1.3 Polak–Ribière

Polak–Ribière is an alternative nonlinear conjugate gradient method, where  $\beta_k$  is computed as follows:

$$\beta_k^{PR} = \frac{\nabla f(x_{k+1})^T (\nabla f(x_{k+1}) - \nabla f(x_k))}{\nabla f(x_k)^T \nabla f(x_k)}$$

### 1.4 Backtracking Line Search

In our case, nonlinear optimization, finding the best step length  $\alpha$  is itself a nonlinear optimization problem. To avoid solving it exactly at every iteration, we use backtracking.

In backtracking we start with  $\alpha = \alpha_0$ , and iteratively shrink it by  $\alpha \leftarrow \rho \alpha$ ,  $\rho \in (0, 1)$ , until the Armijo condition is met:

$$f(x_k + \alpha_k d_k) < f(x_k) + c_1 \alpha \nabla f(x_k)^T d_k$$

## 2 Problem analysis

In our assignment we are tasked with applying the three optimization algorithms to the following function:

$$\min_{x \in \mathbb{R}^n} \sum_{i=1}^n \left( \frac{1}{4} x_i^4 + \frac{1}{2} x_i^2 + x_i \right) = \min_{x \in \mathbb{R}^n} f(x)$$

First of all we compute symbolically the gradient, since we are going to be comparing the exact derivatives with the ones computed via finite differences.

$$(\nabla f(x))_i = x_i^3 + x_i + 1$$

Then, since this problem is additively separable, we can write it as

$$\min_{x \in \mathbb{R}^n} f(x) = \min_{x \in \mathbb{R}^n} \sum_{i=1}^n g(x_i)$$

where  $g(x) = x^4/4 + x^2/2 + x$ . Notice how since  $g$  is convex,  $f$  is convex due to being a sum of  $g$ s.

Since all elements of the sum in  $f$  have the same minimum, then

$$\min_{x \in \mathbb{R}^n} \sum_{i=1}^n g(x_i) = \sum_{i=1}^n \min_{x \in \mathbb{R}^n} g(x_i)$$

Using WolframAlpha we can calculate the minimum for  $g$ , which is  $x^* \approx -0.682328...$ . Thus the optimal  $x \in \mathbb{R}^n$  will be  $(-0.682328, \dots, -0.682328)$ . We will use this to verify the correctness of our algorithms, but we will not exploit separability for minimization, which would allow us to compute the complete solution using just one dimension.

## 2.1 Exploiting separability in the gradient approximation

A naïve implementation of finite differences would require  $O(n)$  evaluations of  $f$ , one for each  $\frac{\partial f(x)}{\partial x_i}$ . In turn, each evaluation of  $f$  requires  $O(n)$  evaluations of  $g$ , because of the summation from 1 to  $n$ .

By exploiting separability we can rewrite the finite differences to require only  $O(n)$  evaluations of  $g$  for the whole gradient, rather than  $O(n^2)$ .

Firstly, we define  $F(x)$ ,  $F : \mathbb{R}^n \rightarrow \mathbb{R}^n$ :

$$F(x) = [g(x_1), g(x_2), \dots, g(x_n)]^T$$

We can get back to our original  $f$  by multiplying with a vector of ones

$$f(x) = \mathbf{1}_n^T F(x)$$

The first finite difference, where  $e_1 = [1, 0, \dots, 0]$  is the first basis vector, then is

$$\begin{aligned} f(x + he_1) - f(x) &= \mathbf{1}_n^T F(x + he_1) - \mathbf{1}_n^T F(x) \\ &= \mathbf{1}_n^T (F(x + he_1) - F(x)) \\ &= \mathbf{1}_n^T \left( \begin{bmatrix} g(x_1 + h) \\ g(x_2) \\ \vdots \\ g(x_n) \end{bmatrix} - \begin{bmatrix} g(x_1) \\ g(x_2) \\ \vdots \\ g(x_n) \end{bmatrix} \right) \\ &= \mathbf{1}_n^T \begin{bmatrix} g(x_1 + h) - g(x_1) \\ 0 \\ \vdots \\ 0 \end{bmatrix} \\ &= g(x_1 + h) - g(x_1) \end{aligned}$$

From this, we can obtain the whole gradient in a single pass with:

$$\nabla^{FD} f(x) = \begin{bmatrix} g(x_1 + h) - g(x_1) \\ g(x_2 + h) - g(x_2) \\ \vdots \\ g(x_n + h) - g(x_n) \end{bmatrix}$$

## 3 Results

### 3.1 Comparing optimization methods

We begin by comparing the computation time and number of iterations for the three optimization methods and the three gradient computation methods. For this comparison we fix the finite difference step size  $h$  to  $-8$ , as well as the number of dimensions  $n$  to  $10^5$ . The computation time is plotted in figure 1 while the number of iterations can be found in figure 2.

We can observe how all three algorithms converge in a similar number of iterations. Steepest descent being on average the fastest with  $\approx 45$  iterations while Polak–Ribière the slowest requiring  $\approx 65$ . Additionally we see that, as expected, the fastest method for calculating the gradient is the exact method. Finite difference methods, despite being optimized to exploit separability, still require double the amount of floating point operations compared to the exact one. This effect is particularly noticeable for Fletcher-Reeves where we see how there is almost a 4x difference in runtime, despite a slightly increase in iterations. Analogous considerations apply for  $n = 10^4$  as seen in tables 1, 2.

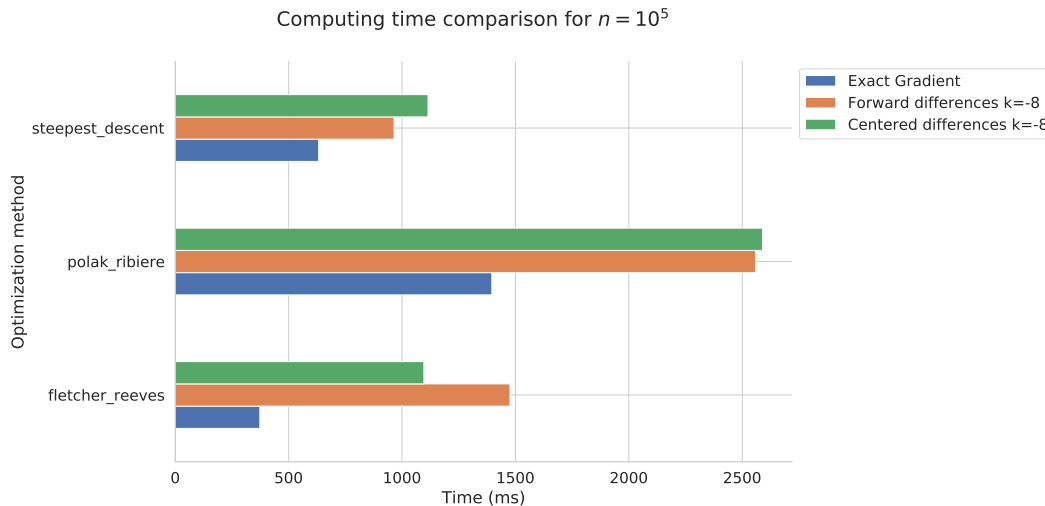


Figure 1: Computing time comparison

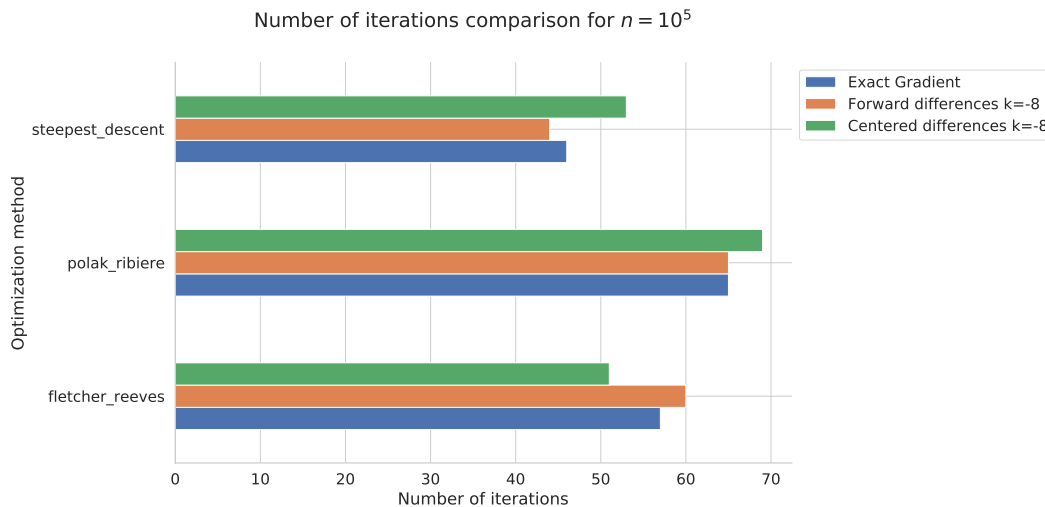


Figure 2: Number of iterations comparison

### 3.2 Effect of finite differences step size

We move on to comparing how different values of  $k$  affect the performance and convergence of the algorithms. The coefficient  $k$  is used to compute the step size  $h = 10^{-k} \|\hat{x}\|$ , where  $\hat{x}$  is the point at which the derivative approximation is computed.

In figures 3 and 4 we see the results for  $n = 10^5$  and using steepest descent. The algorithms fail to converge with step sizes too long i.e.  $k = 2, 4, 6$ , as the number of iterations has reached  $K_{max} = 1000$ . Notice as




Figure 3: Number of iterations for different values of  $k$ , using steepest descent




Figure 4: Computing time

well as for  $k = 6$  the centered differences algorithm maintains enough precision to make steepest descent converge, while forward differences, with their  $O(h)$  error, can not. Similar behaviour is observed for the other optimization methods, as seen in tables 1, 2.

Optimization Method Problem size Gradient method	Computing time					
	Steepest Descent		Fletcher-Reeves		Polak-Ribière	
	4	5	4	5	4	5
Centered k=10	0.104	0.998	0.135	1.318	0.210	2.588
Centered k=12	0.105	0.991	0.113	1.329	0.207	2.811
Centered k=14	0.091	0.983	0.095	1.270	0.222	2.733
Centered k=2	4.271	46.597	2.550	109.914	3.775	49.870
Centered k=4	0.114	43.175	0.078	0.671	0.274	3.196
Centered k=6	0.107	0.985	0.112	1.304	0.210	2.587
Centered k=8	0.109	0.965	0.116	1.476	0.212	2.559
Exact	0.048	0.632	0.047	0.373	0.137	1.396
Forward k=10	0.102	1.014	0.136	1.268	0.204	2.430
Forward k=12	0.106	1.036	0.120	1.499	0.208	2.534
Forward k=14	0.089	0.951	0.091	1.254	0.220	2.866
Forward k=2	3.840	52.684	3.567	114.210	3.902	51.761
Forward k=4	4.264	45.400	2.369	32.173	0.207	1.529
Forward k=6	0.086	46.468	0.073	0.736	0.231	2.913
Forward k=8	0.121	1.114	0.105	1.096	0.204	2.589

Table 1: Computing times (s) for all combinations of methods and parameters

Optimization Method Problem size Gradient method	Iterations					
	Steepest Descent		Fletcher-Reeves		Polak–Ribière	
	4	5	4	5	4	5
Centered k=10	49	47	69	56	65	69
Centered k=12	47	46	58	56	65	71
Centered k=14	43	46	48	50	69	69
Centered k=2	1000	1000	709	1000	1000	1000
Centered k=4	51	1000	38	27	87	87
Centered k=6	48	46	58	56	65	69
Centered k=8	45	44	60	60	65	65
Exact	47	46	69	57	65	65
Forward k=10	48	49	70	57	65	65
Forward k=12	49	48	63	56	65	67
Forward k=14	43	41	45	47	70	79
Forward k=2	1000	1000	1000	1000	1000	1000
Forward k=4	1000	1000	602	556	65	41
Forward k=6	42	1000	38	30	73	73
Forward k=8	57	53	55	51	65	69

Table 2: Number of iterations for all combinations of methods and parameters

## 4 Code

```

1  # The function f we need to optimize
2  function f(x)
3      val = 0.0
4      for i = eachindex(x)
5          val += 1/4 * x[i]^4 + 1/2 * x[i]^2 + x[i]
6      end
7      return val
8  end
9
10 # The gradient of f, computed symbolically
11 ∇f(x) = [x[i]^3 + x[i] + 1 for i = eachindex(x)]
12
13 # Compute the optimal step length via backtracking line
14 # search, using Armijo-Goldstein conditions
15 function backtrack(xk, dk, f, ∇fk, α0, c1, ρ, btmax)
16     α = α0
17     bt = 1
18     fk = f(xk)
19     ∇f_dk = ∇fk' * dk
20     while bt < btmax
21         if (f(xk + α * dk) < fk + c1 * α * ∇f_dk)
22             break
23         end
24         α = ρ * α
25         bt += 1
26     end
27     return bt, α
28 end
29

```

```

30 # Nonlinear steepest descent method.
31 function steepest_descent(
32     x0, # starting point
33     f, ∇f, # function handles for the objective and its gradient
34     rel_diff=1e-8, kmax=1000, # solver hyperparameters
35     α0=5.0, c1=1e-4, ρ=0.8, btmax=50 # backtracking hyperparameters
36 )
37 k = 0
38 xk = x0
39
40 while k < kmax
41     ∇fk = ∇f(xk)
42     dk = -∇fk
43
44     bt, αk = backtrack(xk, dk, f, ∇fk, α0, c1, ρ, btmax)
45
46     xk_new = xk + αk * dk
47     # relative movement stopping criterion
48     if norm(xk_new - xk) / norm(xk) < rel_diff
49         break
50     end
51     xk = xk_new
52
53     k += 1
54 end
55
56 return xk, k
57 end
58
59 function fletcher_reeves(
60     x0,
61     f, ∇f,
62     rel_diff=1e-8, kmax=1000,
63     α0=5.0, c1=1e-4, ρ=0.8, btmax=50
64 )
65 xk = x0
66 ∇fk = ∇f(xk)
67
68 pk = -∇fk
69 k = 0
70
71 while k < kmax
72     bt, αk = backtrack(xk, pk, f, ∇fk, α0, c1, ρ, btmax)
73
74     xkp1 = xk + αk * pk
75     if norm(xkp1 - xk) / norm(xkp1) < rel_diff
76         break
77     end
78
79     ∇fkp1 = ∇f(xkp1)
80     βkp1 = (∇fkp1' * ∇fkp1) / (∇fk' * ∇fk)
81
82     pk = -∇fkp1 + βkp1 * pk
83

```

```

84         xk = xkp1
85         ∇fk = ∇fkp1
86
87         k = k + 1
88     end
89
90     return xk, k
91 end
92
93 function polak_ribiere(
94     x0,
95     f, ∇f,
96     rel_diff=1e-8, kmax=1000,
97     α0=5.0, c1=1e-4, ρ=0.8, btmax=50
98 )
99     xk = x0
100    ∇fk = ∇f(xk)
101
102    pk = -∇fk
103    k = 0
104
105    while k < kmax
106        bt, αk = backtrack(xk, pk, f, ∇fk, α0, c1, ρ, btmax)
107
108        xkp1 = xk + αk * pk
109        if norm(xkp1 - xk) / norm(xkp1) < rel_diff
110            break
111        end
112
113        ∇fkp1 = ∇f(xkp1)
114        βkp1 = (∇fkp1' * (∇fkp1 - ∇fk)) / (∇fk' * ∇fk)
115
116        pk = -∇fkp1 + βkp1 * pk
117
118        xk = xkp1
119        ∇fk = ∇fkp1
120
121        k = k + 1
122    end
123
124    return xk, k
125 end
126
127 g(x) = 1/4 * x^4 + 1/2 * x ^ 2 + x
128
129 function ∇f_fwd_diff(x; k=8)
130     h = 10.0^(-k) * norm(x)
131     return [(g(x[i] + h) - g(x[i])) / h for i = eachindex(x)]
132 end
133
134 function ∇f_cnt_diff(x; k=8)
135     h = 10.0^(-k) * norm(x)
136     return [(g(x[i] + h) - g(x[i] - h)) / 2h for i = eachindex(x)]
137 end

```